

When Cheap Transformations Vanish: A Diagnostic Study of Shallow Symbolic Solvers on ARC-AGI-2

本 **Writeup** 已就绪，无需替换任何占位符。直接粘贴即可。

公开 notebook: <https://www.kaggle.com/code/luoguoqiang/arc-prize-2026-arc-agi-2-v1-diagnostic-engine>

代码提交 ID: 56199696 (submission.json , 240 题, schema 校验通过)

开源仓库 (MIT-o) : <https://github.com/LuoGuoQiang-web/arc-prize-2026-arc-agi-2-v1>

Media Gallery 请上传两张图: cover.png (封面, 必需) 与 coverage_curve.png 。

Subtitle

Coverage, not search, governs ARC-AGI-2 for fixed-vocabulary solvers: we quantify a transformation-family shift, show coverage collapses to a cliff rather than a slope, and recover +44% relative by spending the second attempt on a different hypothesis class.

Track

Paper Track (the same Writeup is also submitted as the ARC-AGI-2 Solution Writeup; the two tracks share one six-dimension rubric).

ARC-AGI-2: shallow solvers collapse because the cheap transformation families were removed




Figure 1 — Family shift, method scores, coverage collapse, and the shape ablation.

1. Claim

Shallow symbolic solvers—program search over a fixed vocabulary of grid transformations—are a standard ARC baseline, yet they score near zero on ARC-AGI-2. This is not an implementation artifact. We introduce a reproducible transformation-family taxonomy, measure a sharp distribution shift between the two public splits, show that the dominant eval family (50.8%) needs object-level conditionals our vocabulary cannot express, and show by controlled ablation that output-shape inference—the usual suspect—is not the bottleneck: correcting shapes recovers exactly zero points.

2. Setup

One discipline, four search spaces. A candidate is accepted only if it reproduces *every* demonstration exactly; parameters (colour maps, per-property palettes, fill colours) are fitted on the demonstrations, never guessed.

- **A — geometry DSL.** 29 primitives (dihedral transforms, crops, gravities, scaling, tiling, mirror completion, uniform row/column removal, extreme-component crops) composed to depth ≤ 2 : ~318 chains per task. Training coverage 2.7%.
- **B — object-level operators.** Components with properties (size, colour, bbox, border contact, squareness, holes); predicates select objects, actions delete, keep or recolour them. Training coverage 1.9%.
- **C — context-conditioned recolouring.** Cells are described by structural features (own colour, its object's size and size-rank, boundary/border/hole/square flags); the map feature-tuple $\rightarrow$ colour is fitted on the demonstrations. Training coverage 3.8%.
- **D — shape-changing and counting.** Size-changing outputs: crop to an object's bbox, erase-a-colour-then-crop, integer scaling, and count-to- 1×1 maps. Training coverage 1.6% with **56%** precision (counting maps fit coincidentally), so it ranks below every other family and above the priors—it can only replace a guess.

Metric and data. Kaggle's metric: two attempts per test input, exact match on either scores 1, final score is the mean. Data: the official ARC-AGI-2 release (1,000 training, 120 public eval tasks). Nothing is trained and no eval feedback was used for tuning.

3. Diagnosis

3.1 Family taxonomy and distribution shift

We label every demonstration pair by the simplest family that explains it (shape-changing; per-cell colour function; uniform object recolouring; pixel removal; pixel addition; otherwise object-level conditional), and give each task its dominant label ($\geq 80\%$).

Family	Training (1000 tasks)	Public eval (120 tasks)
shape-changing	31.6%	30.8%
object-level conditional (F_same_other)	28.9%	50.8%
uniform object recolouring	22.6%	4.2%
per-cell colour function	1.8%	0.0%
pure removal	1.1%	0.0%

Structural probes agree: outputs that are integer tilings of the input fall from 0.3% to **0.0%**, contiguous sub-grids from 5.2% to **0.0%**, tiny counting outputs (≤ 4 cells) from 2.1% to **0.0%**. Eval keeps the hard

families and drops precisely those our vocabulary covers cheaply.

3.2 Coverage collapse, and why it is a cliff rather than a slope




Figure 2 — Coverage curve: graded failure on training, identically zero at every depth on eval.

Coverage—the fraction of tasks admitting at least one verified program—collapses to zero:

	training	public eval
program reproducing every demonstration	27 / 1000	0 / 120
program surviving leave-one-out refitting on ≥ 1 fold	26 / 1000	0 / 120

Binary coverage hides the shape of the failure. Refitting each chain on $k-1$ demonstrations and testing the held-out one yields a *coverage curve*. On training it decays gradually (2.6% $\rightarrow$ 2.4% $\rightarrow$ 1.8% $\rightarrow$ 0.4% $\rightarrow$ 0%), so the vocabulary sits near the boundary for many tasks. On eval the curve is **identically zero at every depth** (Figure 2): across 37,951 chain–task pairs, not one chain generalises to even a single held-out demonstration on any of the 120 tasks. Nor is this one design's artefact: all five search spaces of §2—structurally unrelated, with training coverage 1.6–4.3%—validate on **exactly 0/120** eval tasks. The vocabulary is not slightly too small for eval; it is orthogonal to it.

Correction (2026-09-13): an earlier draft reported 15/240 union coverage on the organisers' test split and an expectation near 4.7%. That was wrong — the shipped `arc-agi_test_challenges.json` is a **placeholder** (all 240 ids are byte-identical copies of training tasks; the real file appears only at rerun). The public eval, where coverage is **0/120**, remains our only honest signal.

3.3 Shape inference is not the bottleneck

Shape prediction is the usual diagnosis (74% of eval failures). We tested it: "output equals input shape" matches **71.3%** of eval inputs, non-background bounding box 56.9%, our fallback's modal-shape rule only 25.1%; yet cropping/padding predictions to the oracle shape yields **0/167** and resizing to it also **0/167**, with 43/167 inputs already correctly shaped but wrong in content. Shape error is a *symptom* of falling back to priors, not an independent defect.

3.4 A floor of trivial baselines

Demonstration-only baselines score 0.09–1.21% on training (identity, modal output, retrieval, per-cell majority); A beats them all, yet **every shallow method is exactly 0.00% on eval**.

4. Intervention: spending the second attempt on a different hypothesis class

ARC-AGI-2 grants exactly two attempts per test input. Let A have hit set H_A and B have H_B . Spending both attempts inside A realises at most H_A ; spending the second on B realises $H_A \cup (H_B \setminus H_A)$. **The return on the second slot is therefore exactly the asymmetric difference $|H_B \setminus H_A|$ —not B's accuracy.** A weaker but disjoint solver can beat a stronger but redundant one.

Configuration (each switchable in the released engine)	Training hits	Score
geometry DSL only	34 / 1076	3.160%
+ object-level family	45 / 1076	4.182%
+ context-conditioned family	49 / 1076	4.554%
+ shape-changing / counting (default)	52 / 1076	4.833%

The first attempt alone scores 36/1076 (3.35%); giving it a second guess from the same solver adds 0.19 pp, whereas the cross-family slots add 1.67 pp—a **9× larger return per slot**. Gains are localised, not diffuse:

Family	n	geometry	+object	+context	+shape/counting
shape-changing	353	18	18	18	21
object-level conditional	314	7	8	11	11
uniform object recolouring	232	1	11	12	12
other (per-cell / add / remove / mixed)	177	8	8	8	8
total	1076	34	45	49	52

Every gain lands exactly where the added family's competence lies—the object families move the two object families, and the shape/counting family moves shape-changing $18 \rightarrow 21$ —with **no family ever regressing**, as the set-difference reading predicts. The transferable rule for anyone working under a fixed attempt budget: measure your solvers' hit sets, then spend the second slot to maximise the difference—not to sample the same model twice.

5. Interpretation: why coverage, not search, governs this distribution

Fix a vocabulary V and a depth budget k ; let $\text{Programs}_{\leq k}(V)$ be what it can express. Define coverage $C(V) = P_{t \sim D}[\exists p \in \text{Programs}_{\leq k}(V) \text{ that reproduces every demonstration of } t]$. Score is bounded by coverage: a test input whose demonstrations admit no expressible program can only be answered by luck. We argue $C(V)$ is governed by *family alignment*, through two mechanisms.

Mechanism 1—family removal. A finite vocabulary induces a set of expressible transformation families $F(V)$. When the evaluation distribution is constructed to exclude the families that $F(V)$ covers cheaply, those primitives become dead weight and coverage falls to zero rather than merely degrading. Per-cell recolouring, tiling, sub-grid extraction and counting are 1.8–5.2% of training and **0.0%** of eval, where coverage is **0/120**—not low, but zero.

Mechanism 2—compositional explosion. Object-conditional tasks are described by a predicate over objects, an action, and optionally a property map. Expressing them requires the product of these choices,

so a hand-built vocabulary grows sub-linearly against a space that grows multiplicatively—which is why the dominant eval family (50.8%) is precisely the one our primitives cannot reach.

The hypothesis is falsifiable, and three predictions hold. (P1) Extending the vocabulary raises hits only inside the extended family's competence: the object family moves uniform recolouring $1 \rightarrow 11$ and object-conditional $7 \rightarrow 8$, the context family lifts object-conditional to 11, and the shape/counting family moves shape-changing $18 \rightarrow 21$, while the other 177 inputs never move. (P2) More search cannot substitute for vocabulary: ~318 chains already validate nothing on eval. (P3) Heterogeneous portfolios pay off by exactly $|H_B \setminus H_A|$: measured +11 hits, that bound. The practical consequence: instrument *coverage per family* before scaling search or model size—a cheap measurement that predicts a symbolic solver's ceiling far better than its training accuracy.

6. Limitations and negative result

We report no improvement on the public evaluation split: **0.000% (0/167)** for A, B, C, their combinations, and every trivial baseline—eval coverage stays **0/120**, so portfolio gains cannot transfer. This is a capability boundary, not a tuning failure: the families our vocabularies cover cheaply are absent from eval by construction. We used the eval split once; our claims concern *mechanism*, not accuracy.

7. Reproducibility

Runs on CPU in seconds. The submission comes from a Kaggle notebook whose embedded engine is byte-checked against the repository source at build time; 39 self-tests, a full dry-run, and an environment-faithful simulation of the submission path (which catches Kaggle-only branches) all pass; one script re-verifies every headline number. Data: official ARC-AGI-2 release (Apache-2.0); code MIT-O.

Links. Code (MIT-O): github.com/LuoGuoQiang-web/arc-prize-2026-arc-agi-2-v1. Notebook: [luoguoqiang/arc-prize-2026-arc-agi-2-v1-diagnostic-engine](#) on Kaggle. Linked code submission: **56199696** ([submission.json](#), 240 tasks, schema-valid).

Figure 1 (cover) shows the four headline measurements; **Figure 2** details the coverage curve.

Appendix: Links

- Code (MIT-O): <https://github.com/LuoGuoQiang-web/arc-prize-2026-arc-agi-2-v1>
- Kaggle notebook (public): <https://www.kaggle.com/code/luoguoqiang/arc-prize-2026-arc-agi-2-v1-diagnostic-engine>
- Linked Kaggle code submission: **56199696** ([submission.json](#), 240 tasks, schema-valid)
- Figures: [cover.png](#) (Figure 1, cover image — required), [coverage_curve.png](#) (Figure 2)
- Data: official ARC-AGI-2 release, Apache-2.0.
- Reproduction: `python work/arc_w1/reproduce_all.py` re-verifies every number above (18/18 checks).

Appendix: Author background

Solo entrant with a background in applied machine learning and software engineering. This project was run as a deliberately bounded study: no GPU cluster, no pretrained weights, no network at inference—only the official ARC-AGI-2 data, a CPU, and a fixed time budget. The goal was to find out *why* a classical program-search baseline collapses on ARC-AGI-2 before attempting to scale anything up.